

파이썬을 활용한 중고차 가격 예측

20181111 박초연

I. 개요

- A. Data Set 확인
- B. 데이터 전처리
- C. 데이터 분석
- D. 분석 결과
- E. 코드 링크

II. Data Set 확인

A. 데이터 불러오기

1. 사용한 데이터는 kaggle의 'Used Cars Price Prediction' 중 train-data 사용.
2. Url : <https://www.kaggle.com/avikasliwal/used-cars-price-prediction>

B. 데이터 확인

| | Un na me d: 0 | Name | Lo ca ti on | Year | Kilom eters _Driv en | Fuel_ Type | Trans missi on | Owner _Type | Mil eag e | Eng ine | Powe r | Seats | New _Pr ice | Price |
|---|---------------------------|----------------------------------|----------------------|------|-------------------------------|---------------|----------------------|----------------|-----------------|------------|-----------|-------|-------------------|-------|
| 0 | 0 | Maruti Wag on R LXI CNG | M u m b a i | 2010 | 72000 | CNG | Manual | First | 26.6 km/kg | 998 CC | 58.16 bhp | 5.0 | Na N | 1.75 |
| 1 | 1 | Hyundai Creta 1.6 CRDi SX Option | Pu ne | 2015 | 41000 | Diesel | Manual | First | 19.67 kmpl | 1582 CC | 126.2 bhp | 5.0 | Na N | 12.50 |

III. 데이터 전처리

A. Column 설명

1. Unnamed: 0 -> index
2. Name -> 자동차 모델명
3. Location -> 자동차의 사용 지역
4. Year -> 출고년도
5. Kilometers_Driven -> 자동차의 주행거리
6. Puel_Type -> 사용되는 연료
7. Transmission -> 변속기 종류
8. Owner_Type -> 현 사용자가 몇번째 사용자인지
9. Mileage -> 연비
10. Engine -> 엔진
11. Power -> 마력
12. Seats -> 좌석 수
13. New_price -> 출고가
14. Price -> 중고가

B. 무의미한 변수 제거

1. 'Unnamed: 0' 안의 값은 인덱스와 같으므로 분석에 필요하지 않다고 판단하여 제외함.

2. 'Name', 'Location'의 경우 데이터 값이 매우 다양함. 이를 통해 어떠한 특징을 이끌어내어, 중고차 가격을 예측하기 어렵다고 판단하여 제외함.
 3. 'New_Price' 데이터의 값을 보면 결측값이 전체의 약 86.31%로 너무 많아 제외함.
- C. 문자형 자료를 숫자형 자료로 변환
1. 'Mileage', 'Engine', 'Power'의 경우 숫자 정보를 가지고 있지만, 뒤에 단위가 붙어 문자형 데이터로 읽어짐. 이러한 문제를 해결하기 위해 단위를 지우고 데이터의 형태를 'float'으로 바꿈.
 2. 'Power'는 공백이 포함되어 단위를 삭제해도 데이터의 형태가 안바뀜. 따라서 공백을 제거함. 또한, 결측값이 'null'이라는 문자로 들어가 있어 NaN 값으로 바꿈.
- D. 범주형 자료를 숫자형 자료로 변환
1. 'Fuel_Type', 'Transmission', 'Owner_Type'은 문자로 되어 있지만, 'Price'값에 영향을 줄 수 있다고 판단하여 숫자로 변경함.
 2. 같은 value를 같은 숫자로 표현함.
- E. 결측값 제거
1. 'Mileage', 'Engine', 'Power', 'Seats' 안에 결측값이 포함되어 있어 결측값을 제거함.
- F. 'Price'의 범주형 자료 만들기
1. 'Price'는 연속형 자료이고 연속형 자료는 예측의 정확도가 떨어지기 때문에 범주를 나눠주면 정확도가 올라감.
 2. 데이터의 비율이 비슷하게 범주를 나누고 'Class' 변수를 만들어 입력함.
 3. 범주는 사분위수 사용함. 75% 이상, 50% 이상, 25% 이상, 25% 이하

| | Year | Kilometers_Driven | Fuel_Type | Transmission | Owner_Type | Mileage | Engine | Power | Seats | Price | Class |
|---|------|-------------------|-----------|--------------|------------|---------|--------|--------|-------|-------|-------|
| 0 | 2010 | 72000 | 3 | 0 | 1 | 26.60 | 998.0 | 58.16 | 5.0 | 1.75 | D |
| 1 | 2015 | 41000 | 1 | 0 | 1 | 19.67 | 1582.0 | 126.20 | 5.0 | 12.50 | A |

IV. 데이터 분석

- A. 다중 선형 회귀 분석
1. 다중 선형 회귀 분석을 통해 'Price'에 영향을 주는 가장 유의미한 변수 식을 도출함.
 2. $Price = 1.0376 * Year - 1.8900 * Fuel_Type + 2.6342 * Transmission - 0.1327 * Mileage + 0.0012 * Engine + 0.1249 * Power - 0.9559 * Seats$
 3. 여러차례 OLS Regression Results 확인 후, 'Kilometers_Driven'과 'Owner_Type'은 'Price'의 변화에 큰 영향을 주지 못한다고 판단함.
- B. 분석에 활용되는 Train & Test data를 변하기 않게 설정
1. shuffle=False
 2. random_state=20181111
 3. random_state=0
- C. 서포트 벡터 머신(SVM) <random_state=20181111>
1. 'linear'
 - a) C=1, gamma=1 : accuracy = 0.7343550446998723
 - b) parameter 최적화 후 C=500, gamma=0.01 : accuracy = 0.7318007662835249
 - c) parameter 최적화 후 C=500, gamma=0.001 : accuracy = 0.7318007662835249
 - d) parameter 최적화 후 C=500, gamma=0.0001 : accuracy = 0.7318007662835249
 - e) GridSearchCV의 gamma 값을 바꿔도 accuracy가 같음.
 2. 'rbf'
 - a) C=1, gamma=1 : accuracy = 0.7692635163899532
 - b) parameter 최적화 후 C=1, gamma=1 : accuracy = 0.7692635163899532

D. 신경망(MLP) <random_state=20181111>

1. 'lbfgs'

- a) hidden_layer_sizes=(10,10) : accuracy = 0.7432950191570882
- b) parameter 최적화 후 alpha =0.1, hidden_layer_sizes=(500,) : accuracy = 0.7420178799489144
- c) parameter 최적화 후 alpha = 0.01, hidden_layer_sizes=(500,500) : accuracy = 0.7420178799489144
- d) GridSearchCV의 hidden_layer_sizes 값을 바꿔도 accuracy가 같음.

E. 서포트 벡터 머신(SVM) <random_state=0>

1. 'linear'

- a) C=1, gamma=1 : accuracy = 0.7343550446998723
- b) parameter 최적화 후 C=500, gamma=0.01 : accuracy = 0.7318007662835249
- c) parameter 최적화 후 C=500, gamma=0.001 : accuracy = 0.7318007662835249
- d) parameter 최적화 후 C=500, gamma=0.0001 : accuracy = 0.7318007662835249
- e) GridSearchCV의 gamma 값을 바꿔도 accuracy가 같음.

2. 'rbf'

- a) C=1, gamma=1 : accuracy = 0.7692635163899532
- b) parameter 최적화 후 C=1, gamma=1 : accuracy = 0.7692635163899532

F. 신경망(MLP) <random_state=0>

1. 'lbfgs'

- a) hidden_layer_sizes=(10,10) : accuracy = 0.7352064708386548
- b) parameter 최적화 후 alpha =0.001, hidden_layer_sizes=(500,500) : accuracy = 0.7432950191570882

V. 분석 결과

- A. 서포트 벡터 머신 중 비선형분류 'rbf' 방식은 Train & Test data가 변하더라도 동일한 parameter와 accuracy가 나왔으며, accuracy가 가장 높게 나옴.
- B. 서포트 벡터 머신 중 비선형분류 'rbf' 방식이 'Price'를 가장 잘 예측함.

VI. 코드 링크

<random_state=20181111>

https://colab.research.google.com/drive/1P0e1DL4_fPNSIAx-jQPNRaPwiCCfmF-F?usp=sharing

<random_state=0>

https://colab.research.google.com/drive/1EUURXWsBNB_oN0AptjA4JzuD-7GXA2yU?usp=sharing

▼ 데이터 불러오기

```
1 from google.colab import drive
2 drive.mount('/content/drive')
```

↳ Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount('/content/drive', force_remount=True).

```
1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
```

```
1 data = pd.read_csv('/content/drive/My Drive/3학년1학기/Python_Application/
2 data.head(2)
```

↳

| | Unnamed: 0 | Name | Location | Year | Kilometers_Driven | Fuel_Type | Transmissi |
|---|------------|----------------------------------|----------|------|-------------------|-----------|------------|
| 0 | 0 | Maruti Wagon R LXI CNG | Mumbai | 2010 | 72000 | CNG | Mar |
| 1 | 1 | Hyundai Creta 1.6 CRDi SX Option | Pune | 2015 | 41000 | Diesel | Mar |

```
1 data.info()
```

```
[>] <class 'pandas.core.frame.DataFrame'>
RangeIndex: 6019 entries, 0 to 6018
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0            6019 non-null  int64
1   Name                  6019 non-null  object
2   Location              6019 non-null  object
3   Year                  6019 non-null  int64
4   Kilometers_Driven     6019 non-null  int64
5   Fuel_Type             6019 non-null  object
6   Transmission          6019 non-null  object
7   Owner_Type            6019 non-null  object
8   Mileage               6017 non-null  object
9   Engine                5983 non-null  object
10  Power                 5983 non-null  object
11  Seats                 5977 non-null  float64
12  New_Price             824 non-null   object
13  Price                 6019 non-null  float64
dtypes: float64(2), int64(3), object(9)
memory usage: 658.5+ KB
```

▼ 데이터 전처리

▶ 무의미한 변수 제거

'Unnamed: 0' 인덱스와 같음

'Name', 'Location' 매우 다양함. 때문에 어떤 특징을 이끌어내서 중고가와 연결 어려움.

'New_Price' 결측값이 너무 많음. 전체의 약 86.31%

↳ 숨겨진 셀 3개

▼ 문자형 자료를 숫자형 자료로 변환

```

1 df['Engine']=df['Engine'].str.replace('CC','').astype(float)
2 df['Mileage']=df['Mileage'].str.replace('kmpl','')
3 df['Mileage']=df['Mileage'].str.replace('km/kg','').astype(float)
4 df['Power']=df['Power'].str.replace('bhp','')
5 df['Power']=df['Power'].str.split(' ').str[0]
6 df['Power'][df['Power'] == 'null'] = np.NaN
7 df['Power']=df['Power'].astype(float)

```

↳ /usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:6: SettingWithA value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: <https://pandas.pydata.org/pandas-docs>

```
1 df.head(2)
```

↳

| | Year | Kilometers_Driven | Fuel_Type | Transmission | Owner_Type | Mileage | Eng |
|---|------|-------------------|-----------|--------------|------------|---------|-----|
| 0 | 2010 | 72000 | CNG | Manual | First | 26.60 | 9 |
| 1 | 2015 | 41000 | Diesel | Manual | First | 19.67 | 15 |

```
1 df.dtypes
```

↳

```

Year                int64
Kilometers_Driven   int64
Fuel_Type           object
Transmission         object
Owner_Type          object
Mileage             float64
Engine              float64
Power               float64
Seats               float64
Price               float64
dtype: object

```

▼ 범주형 자료를 숫자형 자료로 변환

```
1 df['Fuel_Type'].value_counts()
```

↳

```

Diesel      3205
Petrol      2746
CNG         56
LPG         10
Electric     2
Name: Fuel_Type, dtype: int64

```

```

1 df['Fuel_Type'][df['Fuel_Type']=='Diesel'] = 1
2 df['Fuel_Type'][df['Fuel_Type']=='Petrol'] = 2
3 df['Fuel_Type'][df['Fuel_Type']=='CNG'] = 3
4 df['Fuel_Type'][df['Fuel_Type']=='LPG'] = 4
5 df['Fuel_Type'][df['Fuel_Type']=='Electric'] = 5
6 df['Fuel_Type']=df['Fuel_Type'].astype(int)
7 df['Transmission'][df['Transmission']=='Manual'] = 0
8 df['Transmission'][df['Transmission']=='Automatic'] = 1
9 df['Transmission']=df['Transmission'].astype(int)
10 df['Owner_Type'][df['Owner_Type']=='First'] = 1
11 df['Owner_Type'][df['Owner_Type']=='Second'] = 2
12 df['Owner_Type'][df['Owner_Type']=='Third'] = 3
13 df['Owner_Type'][df['Owner_Type']=='Fourth & Above'] = 4
14 df['Owner_Type']=df['Owner_Type'].astype(int)

```

```
1 df.head(2)
```

```

↳
   Year  Kilometers_Driven  Fuel_Type  Transmission  Owner_Type  Mileage  Engine
0  2010                72000          3             0           1    26.60      9
1  2015                41000          1             0           1    19.67     15

```

```
1 df.dtypes
```

```

↳ Year                int64
   Kilometers_Driven    int64
   Fuel_Type            int64
   Transmission          int64
   Owner_Type            int64
   Mileage              float64
   Engine               float64
   Power                float64
   Seats               float64
   Price               float64
   dtype: object

```

▼ 결측값 제거

```
1 df.info()
```

```
[>] <class 'pandas.core.frame.DataFrame'>
RangeIndex: 6019 entries, 0 to 6018
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Year                  6019 non-null  int64
1   Kilometers_Driven     6019 non-null  int64
2   Fuel_Type             6019 non-null  int64
3   Transmission          6019 non-null  int64
4   Owner_Type            6019 non-null  int64
5   Mileage               6017 non-null  float64
6   Engine                5983 non-null  float64
7   Power                 5876 non-null  float64
8   Seats                 5977 non-null  float64
9   Price                 6019 non-null  float64
dtypes: float64(5), int64(5)
memory usage: 470.4 KB
```

```
1 df1 = df.dropna()
2 df1 = df1.reset_index(drop=True)
3 df1.head(2)
```

```
[>]
```

| | Year | Kilometers_Driven | Fuel_Type | Transmission | Owner_Type | Mileage | Engine |
|---|------|-------------------|-----------|--------------|------------|---------|--------|
| 0 | 2010 | 72000 | 3 | 0 | 1 | 26.60 | 9 |
| 1 | 2015 | 41000 | 1 | 0 | 1 | 19.67 | 15 |

```
1 df1.isnull().sum()
```

```
[>] Year                0
Kilometers_Driven      0
Fuel_Type              0
Transmission           0
Owner_Type             0
Mileage               0
Engine                0
Power                 0
Seats                 0
Price                 0
dtype: int64
```

▼ Price의 범주형 자료 만들기

75% 이상, 50% 이상, 25% 이상, 25% 이하

연속형 자료는 예측의 정확도가 떨어짐.


```
1 df1.Price.describe()
```

```
count    5872.000000
mean       9.603919
std       11.249453
min        0.440000
25%        3.517500
50%        5.750000
75%       10.000000
max       160.000000
Name: Price, dtype: float64
```

```
1 n = df1.shape[0]
2 label = pd.DataFrame(index=df1.index, columns=['Class'])
3 for i in range(5872):
4     j = df1.Price.loc[i]
5     if j >= 10.000000:
6         label.loc[i, 'Class'] = "A"
7     elif j >= 5.750000:
8         label.loc[i, 'Class'] = "B"
9     elif j >= 3.517500:
10        label.loc[i, 'Class'] = "C"
11    else:
12        label.loc[i, 'Class'] = "D"
```

```
1 label['Class'].value_counts()
```

```
A      1485
D      1468
C      1466
B      1453
Name: Class, dtype: int64
```


```
1 df2 = pd.merge(df1, label, left_on=df1.index, right_on=label.index).iloc[:,
2 df2.head(2)
```

```
Year  Kilometers_Driven  Fuel_Type  Transmission  Owner_Type  Mileage  Enc
0  2010                  72000        3             0           1      26.60    9
1  2015                  41000        1             0           1      19.67   15
```

▼ 다중 선형 회귀분석

1 # 산점도 행렬

2 p = pd.plotting.scatter_matrix(df2, figsize=(8, 8), marker='o')



```

1 from statsmodels.formula.api import ols
2 reg_simp = ols('Price ~ Year + Fuel_Type + Transmission + Mileage + Engi
3 reg_simp.summary()

```



OLS Regression Results

| | | | |
|--------------------------|------------------|----------------------------|-----------|
| Dep. Variable: | Price | R-squared: | 0.695 |
| Model: | OLS | Adj. R-squared: | 0.695 |
| Method: | Least Squares | F-statistic: | 1912. |
| Date: | Tue, 07 Jul 2020 | Prob (F-statistic): | 0.00 |
| Time: | 13:23:31 | Log-Likelihood: | -19054. |
| No. Observations: | 5872 | AIC: | 3.812e+04 |
| Df Residuals: | 5864 | BIC: | 3.818e+04 |
| Df Model: | 7 | | |

Covariance Type: nonrobust

| | coef | std err | t | P> t | [0.025 | 0.975] |
|---------------------|------------|---------|---------|-------|-----------|-----------|
| Intercept | -2086.1610 | 56.300 | -37.054 | 0.000 | -2196.530 | -1975.792 |
| Year | 1.0376 | 0.028 | 36.902 | 0.000 | 0.982 | 1.093 |
| Fuel_Type | -1.8900 | 0.201 | -9.408 | 0.000 | -2.284 | -1.496 |
| Transmission | 2.6342 | 0.243 | 10.850 | 0.000 | 2.158 | 3.110 |
| Mileage | -0.1327 | 0.030 | -4.448 | 0.000 | -0.191 | -0.074 |
| Engine | 0.0012 | 0.000 | 2.958 | 0.003 | 0.000 | 0.002 |
| Power | 0.1249 | 0.004 | 31.797 | 0.000 | 0.117 | 0.133 |
| Seats | -0.9559 | 0.137 | -6.984 | 0.000 | -1.224 | -0.688 |

Omnibus: 4452.126 **Durbin-Watson:** 2.039

Prob(Omnibus): 0.000 **Jarque-Bera (JB):** 359805.135

Skew: 2.983 **Prob(JB):** 0.00

Kurtosis: 40.882 **Cond. No.** 1.82e+06

Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 1.82e+06. This might indicate that there are strong multicollinearity or other numerical problems.

▼ Train & Test data

```

1 from sklearn.preprocessing import StandardScaler
2 from sklearn.model_selection import train_test_split

```

```

1 ### 입력변수(x)와 출력변수(y)
2 x = df2[['Year', 'Fuel_Type', 'Transmission', 'Mileage', 'Engine', 'Power', 'S
3 x_std = StandardScaler().fit_transform(x)
4 y = df2.Class

```

```

1 ### Train & Test data
2 x_train, x_test, y_train, y_test = train_test_split(x_std,y,test_size=0.
3                                                    random_state=2018111

```

▼ 서포트 벡터 머신(SVM)

```

1 from sklearn.svm import SVC

```

▼ 선형분류(linear)

```

1 ### SVM
2 svc = SVC(kernel='linear', C=1, gamma=1)
3 model = svc.fit(x_train, y_train)

```

```

1 ### 예측
2 y_pred = model.predict(x_test)
3 y_pred

```

```

[> array(['B', 'C', 'C', ..., 'B', 'D', 'D'], dtype=object)

```

```

1 ### accuracy
2 model.score(x_test, y_test)

```

```

[> 0.7343550446998723

```

```

1 ### 교차표
2 pd.crosstab(y_test, y_pred)

```

```

[>
  col_0  A    B    C    D
  Class
  A    498   69    3    0
  B     54  404  131    3
  C      8  114  373  108
  D      1   14  119  450

```

```

1 from sklearn.model_selection import GridSearchCV
2 # Set the parameters by cross-validation
3 tuned_parameters = {'kernel': ['linear'],
4                      'C': [0.01, 0.1, 1, 10, 50, 100, 500, 1000],
5                      'gamma': [0.01, 0.05, 0.1, 0.5, 1, 10, 50]}

```

```

1 grid = GridSearchCV(SVC(), tuned_parameters)
2 %time grid.fit(x_train, y_train)

```

```

☞ CPU times: user 12min 23s, sys: 116 ms, total: 12min 23s
Wall time: 12min 24s
GridSearchCV(cv=None, error_score=nan,
             estimator=SVC(C=1.0, break_ties=False, cache_size=200,
                           class_weight=None, coef0=0.0,
                           decision_function_shape='ovr', degree=3,
                           gamma='scale', kernel='rbf', max_iter=-1,
                           probability=False, random_state=None, shrinking=
                           tol=0.001, verbose=False),
             iid='deprecated', n_jobs=None,
             param_grid={'C': [0.01, 0.1, 1, 10, 50, 100, 500, 1000],
                         'gamma': [0.001, 0.01, 0.05, 0.1, 0.5, 1, 10, 50],
                         'kernel': ['linear']}},
             pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
             scoring=None, verbose=0)

```

```

1 grid.best_params_

```

```

☞ {'C': 500, 'gamma': 0.001, 'kernel': 'linear'}

```

```

1 ### SVM
2 # svc = SVC(kernel="linear",C=500, gamma=0.01)
3 svc = SVC(kernel="linear",C=grid.best_params_['C'], gamma=grid.best_params_['gamma'])
4 model = svc.fit(x_train, y_train)

```

```

1 ### 예측
2 y_pred = model.predict(x_test)

```

```

1 ### accuracy
2 model.score(x_test, y_test)

```

```

☞ 0.7318007662835249

```

```
1 ### 교차표
2 pd.crosstab(y_test, y_pred)
```

☞

| col_0 | A | B | C | D |
|-------|-----|-----|-----|-----|
| Class | | | | |
| A | 495 | 73 | 2 | 0 |
| B | 54 | 401 | 134 | 3 |
| C | 8 | 112 | 377 | 106 |
| D | 1 | 14 | 123 | 446 |

▼ 비선형분류(rbf)

```
1 ### SVM
2 svc = SVC(kernel="rbf",C=1, gamma=1)
3 model = svc.fit(x_train, y_train)
```

```
1 ### 예측
2 y_pred = model.predict(x_test)
```

```
1 ### accuracy
2 model.score(x_test, y_test)
```

☞ 0.7692635163899532

```
1 ### 교차표
2 pd.crosstab(y_test, y_pred)
```

☞

| col_0 | A | B | C | D |
|-------|-----|-----|-----|-----|
| Class | | | | |
| A | 519 | 49 | 2 | 0 |
| B | 52 | 402 | 135 | 3 |
| C | 8 | 81 | 434 | 80 |
| D | 5 | 12 | 115 | 452 |

```

1 from sklearn.model_selection import GridSearchCV
2 tuned_parameters = {'kernel': ['rbf'],
3                       'C': [0.01, 0.1, 1, 10, 50],
4                       'gamma': [0.1, 0.5, 1, 10, 50]}

```

```

1 grid = GridSearchCV(SVC(), tuned_parameters)
2 %time grid.fit(x_train, y_train)

```

```

☞ CPU times: user 53.4 s, sys: 15 ms, total: 53.4 s
Wall time: 53.5 s
GridSearchCV(cv=None, error_score=nan,
             estimator=SVC(C=1.0, break_ties=False, cache_size=200,
                           class_weight=None, coef0=0.0,
                           decision_function_shape='ovr', degree=3,
                           gamma='scale', kernel='rbf', max_iter=-1,
                           probability=False, random_state=None, shrinking=
                           tol=0.001, verbose=False),
             iid='deprecated', n_jobs=None,
             param_grid={'C': [0.01, 0.1, 1, 10, 50],
                         'gamma': [0.1, 0.5, 1, 10, 50], 'kernel': ['rbf']}
             pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
             scoring=None, verbose=0)

```

```

1 grid.best_params_

```

```

☞ {'C': 1, 'gamma': 1, 'kernel': 'rbf'}

```

```

1 ### SVM
2 # svc = SVC(kernel="rbf",C=10, gamma=0.1)
3 svc = SVC(kernel="rbf",C=grid.best_params_['C'], gamma=grid.best_params_
4 model = svc.fit(x_train, y_train)

```

```

1 ### 예측
2 y_pred = model.predict(x_test)

```

```

1 ### accuracy
2 model.score(x_test, y_test)

```

```

☞ 0.7692635163899532

```

```
1 ### 교차표
2 pd.crosstab(y_test, y_pred)
```

```
↳
```

| col_0 | A | B | C | D |
|-------|-----|-----|-----|-----|
| Class | | | | |
| A | 519 | 49 | 2 | 0 |
| B | 52 | 402 | 135 | 3 |
| C | 8 | 81 | 434 | 80 |
| D | 5 | 12 | 115 | 452 |

▼ 신경망(MLP)

```
1 from sklearn.neural_network import MLPClassifier
```

```
1 print(x_std.mean(axis=0), x_std.var(axis=0))
```

```
↳ [ 2.29081059e-14  7.62332976e-17 -3.63015703e-17 -4.17468058e-17
    -1.08904711e-17 -2.00868689e-16  5.25162717e-16] [1. 1. 1. 1. 1. 1. 1.]
```

```
1 ### [0,1] 조정
2 x_train, x_test, y_train, y_test = train_test_split(x,y,test_size=0.4)
3 # x_train
4 x_train_min = x_train.min(axis=0)
5 x_train_max = (x_train - x_train_min).max(axis=0)
6 x_train = (x_train - x_train_min) / x_train_max
7 # x_test
8 x_test = (x_test - x_train_min) / x_train_max
```

```
1 ### MLP
2 mlp = MLPClassifier(solver='lbfgs', hidden_layer_sizes=(10,10))
3 model = mlp.fit(x_train, y_train)
```

```
↳ /usr/local/lib/python3.6/dist-packages/sklearn/neural_network/_multilayer_p
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
self.n_iter_ = _check_optimize_result("lbfgs", opt_res, self.max_iter)


```
1 ### 예측
2 y_pred = model.predict(x_test)
```

```
1 ### accuracy
2 model.score(x_test, y_test)
```

```
☞ 0.7432950191570882
```

```
1 ### 교차표
2 pd.crosstab(y_test, y_pred)
```

```
☞
```

| col_0 | A | B | C | D |
|-------|-----|-----|-----|-----|
| Class | | | | |
| A | 551 | 49 | 3 | 0 |
| B | 55 | 387 | 76 | 3 |
| C | 8 | 157 | 345 | 109 |
| D | 1 | 14 | 128 | 463 |

```
1 from sklearn.model_selection import GridSearchCV
2 tuned_parameters = {'solver':['lbfgs'],
3                     'alpha':[0.01,0.1,1,10],
4                     'hidden_layer_sizes':[(5,5),(10,10),(100,),(100,100),
```

```
1 grid = GridSearchCV(MLPClassifier(),tuned_parameters)
2 %time grid.fit(x_train, y_train)
```

```
1 grid.best_params_
```

```
☞ {'alpha': 0.01, 'hidden_layer_sizes': (500, 500), 'solver': 'lbfgs'}
```

```

1 ### MLP
2 mlp = MLPClassifier(solver='lbfgs', alpha=grid.best_params_['alpha'],
3                     hidden_layer_sizes=grid.best_params_['hidden_layer_s
4 model = mlp.fit(x_train, y_train)

```

```

[> /usr/local/lib/python3.6/dist-packages/sklearn/neural_network/_multilayer_p
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

```
self.n_iter_ = _check_optimize_result("lbfgs", opt_res, self.max_iter)
```

```

1 ### 예측
2 y_pred = model.predict(x_test)

```

```

1 ### accuracy
2 model.score(x_test, y_test)

```

```
[> 0.7420178799489144
```

```

1 ### 교차표
2 pd.crosstab(y_test, y_pred)

```

```
[>
```

| col_0 | A | B | C | D |
|--------------|-----|-----|-----|-----|
| Class | | | | |
| A | 548 | 53 | 2 | 0 |
| B | 59 | 371 | 89 | 2 |
| C | 9 | 148 | 353 | 109 |
| D | 1 | 13 | 121 | 471 |

▶ 데이터 불러오기

↳ 숨겨진 셀 4개

▶ 데이터 전처리

↳ 숨겨진 셀 22개

▶ 다중 선형 회귀분석

↳ 숨겨진 셀 2개

▼ Train & Test data

```
1 from sklearn.preprocessing import StandardScaler
2 from sklearn.model_selection import train_test_split

1 ### 입력변수(x)와 출력변수(y)
2 x = df2[['Year', 'Fuel_Type', 'Transmission', 'Mileage', 'Engine', 'Power', 'S
3 x_std = StandardScaler().fit_transform(x)
4 y = df2.Class

1 ### Train & Test data
2 x_train, x_test, y_train, y_test = train_test_split(x_std, y, test_size=0.
3                                     random_state=0)
```

▼ 서포트 벡터 머신(SVM)

```
1 from sklearn.svm import SVC
```

▼ 선형분류(linear)

```

1 ### SVM
2 svc = SVC(kernel='linear', C=1, gamma=1)
3 model = svc.fit(x_train, y_train)

```

```

1 ### 예측
2 y_pred = model.predict(x_test)
3 y_pred

```

```

[> array(['B', 'C', 'C', ..., 'B', 'D', 'D'], dtype=object)

```

```

1 ### accuracy
2 model.score(x_test, y_test)

```

```

[> 0.7343550446998723

```

```

1 ### 교차표
2 pd.crosstab(y_test, y_pred)

```

```

[>
col_0  A    B    C    D
Class
A      498   69    3    0
B       54  404  131    3
C        8  114  373  108
D         1   14  119  450

```

```

1 from sklearn.model_selection import GridSearchCV
2 # Set the parameters by cross-validation
3 tuned_parameters = {'kernel': ['linear'],
4                      'C': [0.1, 1, 10, 50, 100, 500, 1000],
5                      'gamma': [0.01, 0.05, 0.1, 0.5, 1, 10]}

```

```
1 grid = GridSearchCV(SVC(), tuned_parameters)
2 %time grid.fit(x_train, y_train)
```

```
[> CPU times: user 9min 44s, sys: 67.8 ms, total: 9min 44s
Wall time: 9min 44s
GridSearchCV(cv=None, error_score=nan,
             estimator=SVC(C=1.0, break_ties=False, cache_size=200,
                           class_weight=None, coef0=0.0,
                           decision_function_shape='ovr', degree=3,
                           gamma='scale', kernel='rbf', max_iter=-1,
                           probability=False, random_state=None, shrinking=
                           tol=0.001, verbose=False),
             iid='deprecated', n_jobs=None,
             param_grid={'C': [0.1, 1, 10, 50, 100, 500, 1000],
                         'gamma': [0.01, 0.05, 0.1, 0.5, 1, 10],
                         'kernel': ['linear']},
             pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
             scoring=None, verbose=0)
```

```
1 grid.best_params_
```

```
[> {'C': 500, 'gamma': 0.01, 'kernel': 'linear'}
```

```
1 ### SVM
2 # svc = SVC(kernel="linear",C=500, gamma=0.01)
3 svc = SVC(kernel="linear",C=grid.best_params_['C'], gamma=grid.best_para
4 model = svc.fit(x_train, y_train)
```

```
1 ### 예측
2 y_pred = model.predict(x_test)
```

```
1 ### accuracy
2 model.score(x_test, y_test)
```

```
[> 0.7318007662835249
```

```
1 ### 교차표
2 pd.crosstab(y_test, y_pred)
```

☞

| col_0 | A | B | C | D |
|-------|-----|-----|-----|-----|
| Class | | | | |
| A | 495 | 73 | 2 | 0 |
| B | 54 | 401 | 134 | 3 |
| C | 8 | 112 | 377 | 106 |
| D | 1 | 14 | 123 | 446 |

▼ 비선형분류(rbf)

```
1 ### SVM
2 svc = SVC(kernel="rbf",C=1, gamma=1)
3 model = svc.fit(x_train, y_train)
```

```
1 ### 예측
2 y_pred = model.predict(x_test)
```

```
1 ### accuracy
2 model.score(x_test, y_test)
```

☞ 0.7692635163899532

```
1 ### 교차표
2 pd.crosstab(y_test, y_pred)
```

☞

| col_0 | A | B | C | D |
|-------|-----|-----|-----|-----|
| Class | | | | |
| A | 519 | 49 | 2 | 0 |
| B | 52 | 402 | 135 | 3 |
| C | 8 | 81 | 434 | 80 |
| D | 5 | 12 | 115 | 452 |

```

1 from sklearn.model_selection import GridSearchCV
2 tuned_parameters = {'kernel': ['rbf'],
3                       'C': [0.01, 0.1, 1, 10, 50],
4                       'gamma': [0.1, 0.5, 1, 10, 50]}

```

```

1 grid = GridSearchCV(SVC(), tuned_parameters)
2 %time grid.fit(x_train, y_train)

```

```

☞ CPU times: user 54.7 s, sys: 18.9 ms, total: 54.8 s
Wall time: 54.8 s
GridSearchCV(cv=None, error_score=nan,
             estimator=SVC(C=1.0, break_ties=False, cache_size=200,
                           class_weight=None, coef0=0.0,
                           decision_function_shape='ovr', degree=3,
                           gamma='scale', kernel='rbf', max_iter=-1,
                           probability=False, random_state=None, shrinking=
                           tol=0.001, verbose=False),
             iid='deprecated', n_jobs=None,
             param_grid={'C': [0.01, 0.1, 1, 10, 50],
                         'gamma': [0.1, 0.5, 1, 10, 50], 'kernel': ['rbf']}
             pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
             scoring=None, verbose=0)

```

```

1 grid.best_params_

```

```

☞ {'C': 1, 'gamma': 1, 'kernel': 'rbf'}

```

```

1 ### SVM
2 # svc = SVC(kernel="rbf",C=10, gamma=0.1)
3 svc = SVC(kernel="rbf",C=grid.best_params_['C'], gamma=grid.best_params_
4 model = svc.fit(x_train, y_train)

```

```

1 ### 예측
2 y_pred = model.predict(x_test)

```

```

1 ### accuracy
2 model.score(x_test, y_test)

```

```

☞ 0.7692635163899532

```

```
1 ### 교차표
2 pd.crosstab(y_test, y_pred)
```

```
↳
```

| col_0 | A | B | C | D |
|-------|-----|-----|-----|-----|
| Class | | | | |
| A | 519 | 49 | 2 | 0 |
| B | 52 | 402 | 135 | 3 |
| C | 8 | 81 | 434 | 80 |
| D | 5 | 12 | 115 | 452 |

▼ 신경망(MLP)

```
1 from sklearn.neural_network import MLPClassifier
```

```
1 print(x_std.mean(axis=0), x_std.var(axis=0))
```

```
↳ [ 2.29081059e-14  7.62332976e-17 -3.63015703e-17 -4.17468058e-17
    -1.08904711e-17 -2.00868689e-16  5.25162717e-16] [1. 1. 1. 1. 1. 1. 1.]
```

```
1 ### [0,1] 조정
2 x_train, x_test, y_train, y_test = train_test_split(x,y,test_size=0.4)
3 # x_train
4 x_train_min = x_train.min(axis=0)
5 x_train_max = (x_train - x_train_min).max(axis=0)
6 x_train = (x_train - x_train_min) / x_train_max
7 # x_test
8 x_test = (x_test - x_train_min) / x_train_max
```

```
1 ### MLP
2 mlp = MLPClassifier(solver='lbfgs', hidden_layer_sizes=(10,10))
3 model = mlp.fit(x_train, y_train)
```

```
↳ /usr/local/lib/python3.6/dist-packages/sklearn/neural_network/_multilayer_p
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
 self.n_iter_ = _check_optimize_result("lbfgs", opt_res, self.max_iter)


```
1 ### 예측
2 y_pred = model.predict(x_test)
```

```
1 ### accuracy
2 model.score(x_test, y_test)
```

```
0.7352064708386548
```

```
1 ### 교차표
2 pd.crosstab(y_test, y_pred)
```

```
col_0  A    B    C    D
Class
A      530   61    2    1
B       64  375  125    4
C       11  122  368   83
D        1   23  125  454
```

```
1 from sklearn.model_selection import GridSearchCV
2 tuned_parameters = {'solver':['lbfgs'],
3                     'alpha':[0.0001,0.001,0.01,0.1,1,10],
4                     'hidden_layer_sizes':[(100,),(100,100),(500,),(500,500)]}
```

```
1 grid = GridSearchCV(MLPClassifier(),tuned_parameters)
2 %time grid.fit(x_train, y_train)
```

```
1 grid.best_params_
```

```
{'alpha': 0.001, 'hidden_layer_sizes': (500, 500), 'solver': 'lbfgs'}
```

```

1 ### MLP
2 mlp = MLPClassifier(solver='lbfgs', alpha=grid.best_params_['alpha'],
3                     hidden_layer_sizes=grid.best_params_['hidden_layer_s
4 model = mlp.fit(x_train, y_train)

```

```

[> /usr/local/lib/python3.6/dist-packages/sklearn/neural_network/_multilayer_p
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
 self.n_iter_ = _check_optimize_result("lbfgs", opt_res, self.max_iter)

```

1 ### 예측
2 y_pred = model.predict(x_test)

```

```

1 ### accuracy
2 model.score(x_test, y_test)

```

```

[> 0.7432950191570882

```

```

1 ### 교차표
2 pd.crosstab(y_test, y_pred)

```

```

[>
  col_0  A    B    C    D
Class
A      537  52    2    3
B       63 381 120    4
C       10 130 368   76
D         1  20 122 460

```